

part2

March 27, 2024

```
[1]: import pandas as pd
```

```
[2]: data = pd.read_csv("data/published_images.csv")
data.head()
```

```
[2]:
```

	uuid \	iiifurl \	iiifthumburl	viewtype	sequence \	width	height	maxpixels	created	modified \	depictstmsobjectid	assistivetext
0	00007f61-4922-417b-8f27-893ea328206c	https://api.nga.gov/iiif/00007f61-4922-417b-8f...	https://api.nga.gov/iiif/00007f61-4922-417b-8f...	primary	0.0	3365	4332	NaN	2013-07-05 15:41:08-04	2023-07-27 12:06:38-04	17387	NaN
1	0000bd8c-39de-4453-b55d-5e28a9beed38	https://api.nga.gov/iiif/0000bd8c-39de-4453-b5...	https://api.nga.gov/iiif/0000bd8c-39de-4453-b5...	primary	0.0	3500	4688	NaN	2013-08-05 14:31:59-04	2023-07-27 12:11:57-04	19245	NaN
2	0001668a-dd1c-48e8-9267-b6d1697d43c8	https://api.nga.gov/iiif/0001668a-dd1c-48e8-92...	https://api.nga.gov/iiif/0001668a-dd1c-48e8-92...	primary	0.0	3446	4448	NaN	2014-01-02 14:50:50-05	2023-07-27 12:39:11-04	23830	NaN
3	00032658-8a7a-44e3-8bb8-df8c172f521d	https://api.nga.gov/iiif/00032658-8a7a-44e3-8b...	https://api.nga.gov/iiif/00032658-8a7a-44e3-8b...	primary	0.0	2674	3798	NaN	2010-10-13 15:37:25-04	2023-07-27 15:51:54-04	713	NaN
4	0003d4e4-d7fd-4835-8d27-1e9e20672e39	https://api.nga.gov/iiif/0003d4e4-d7fd-4835-8d...	https://api.nga.gov/iiif/0003d4e4-d7fd-4835-8d...	primary	0.0	3000	2648	640.0	2014-11-19 14:24:42-05	2023-11-07 14:13:17-05		

4 71457 NaN

```
[3]: objects = pd.read_csv("data/objects.csv")
      objects.head()
```

/tmp/ipykernel_18622/1759643199.py:1: DtypeWarning: Columns (29) have mixed types. Specify dtype option on import or set low_memory=False.

```
objects = pd.read_csv("data/objects.csv")
```

```
[3]:
```

	objectid	accessioned	accessionnum	locationid	\
0	0	1	1937.1.2.c	NaN	
1	1	1	1937.1.3	NaN	
2	17	1	1937.1.15	NaN	
3	70	1	1937.1.63	8469.0	
4	25	1	1937.1.23	8388.0	

		title	displaydate	beginyear	endyear	\
0		Saint James Major	c. 1310	1310.0	1310.0	
1	Saint Paul and a Group of Worshippers		1333	1333.0	1333.0	
2		Matteo Olivieri (?)	1430s	1430.0	1440.0	
3	An Old Woman Dozing over a Book		c. 1655	1655.0	1655.0	
4		Profile Portrait of a Lady	c. 1410	1410.0	1410.0	

	visualbrowsertimespan		medium	\
0	1300 to 1400		tempera on panel	
1	1300 to 1400		tempera on panel	
2	1401 to 1500	tempera (and oil?) on panel	transferred to canvas	
3	1651 to 1700		oil on canvas	
4	1401 to 1500		oil on panel	

	... parentid	isvirtual	departmentabbr	portfolio	series	volume	watermarks	\
0	...	34.0	0	CIS-R	NaN	NaN	NaN	
1	...	NaN	0	CIS-R	NaN	NaN	NaN	
2	...	NaN	0	CIS-R	NaN	NaN	NaN	
3	...	NaN	0	CNE-B	NaN	NaN	NaN	
4	...	NaN	0	CNE-R	NaN	NaN	NaN	

	lastdetectedmodification	wikidataid	customprinturl
0	2023-05-09 17:01:03.48-04	Q20172973	NaN
1	2024-01-26 22:01:48.797-05	Q20173083	NaN
2	2023-09-15 22:01:37.343-04	Q20173485	NaN
3	2023-05-09 17:01:03.48-04	Q20177396	NaN
4	2023-05-09 17:01:03.48-04	Q3937690	NaN

[5 rows x 30 columns]

```
[4]: data = data.merge(objects[["objectid", "visualbrowsertimespan"]],
↳ left_on="depictstmsobjectid", right_on="objectid", how="left")
data.head()
```

```
[4]:
```

	uuid \	iiifurl \	iiifthumburl	viewtype	sequence \	width	height	maxpixels	created	modified \
0	00007f61-4922-417b-8f27-893ea328206c	https://api.nga.gov/iiif/00007f61-4922-417b-8f...	https://api.nga.gov/iiif/00007f61-4922-417b-8f...	primary	0.0	3365	4332	NaN	2013-07-05 15:41:08-04	2023-07-27 12:06:38-04
1	0000bd8c-39de-4453-b55d-5e28a9beed38	https://api.nga.gov/iiif/0000bd8c-39de-4453-b5...	https://api.nga.gov/iiif/0000bd8c-39de-4453-b5...	primary	0.0	3500	4688	NaN	2013-08-05 14:31:59-04	2023-07-27 12:11:57-04
2	0001668a-dd1c-48e8-9267-b6d1697d43c8	https://api.nga.gov/iiif/0001668a-dd1c-48e8-92...	https://api.nga.gov/iiif/0001668a-dd1c-48e8-92...	primary	0.0	3446	4448	NaN	2014-01-02 14:50:50-05	2023-07-27 12:39:11-04
3	00032658-8a7a-44e3-8bb8-df8c172f521d	https://api.nga.gov/iiif/00032658-8a7a-44e3-8b...	https://api.nga.gov/iiif/00032658-8a7a-44e3-8b...	primary	0.0	2674	3798	NaN	2010-10-13 15:37:25-04	2023-07-27 15:51:54-04
4	0003d4e4-d7fd-4835-8d27-1e9e20672e39	https://api.nga.gov/iiif/0003d4e4-d7fd-4835-8d...	https://api.nga.gov/iiif/0003d4e4-d7fd-4835-8d...	primary	0.0	3000	2648	640.0	2014-11-19 14:24:42-05	2023-11-07 14:13:17-05

	depictstmsobjectid	assistivetext	objectid	visualbrowsertimespan
0	17387	NaN	17387.0	1926 to 1950
1	19245	NaN	19245.0	1926 to 1950
2	23830	NaN	23830.0	1926 to 1950
3	713	NaN	713.0	1501 to 1550
4	71457	NaN	71457.0	1976 to 2000

```
[5]: print(f"Number of rows before dropping NaN values: {data.shape[0]}")
data.dropna(subset=['visualbrowsertimespan'], inplace=True)
print(f"Number of rows after dropping NaN values: {data.shape[0]}")
```

Number of rows before dropping NaN values: 116251
Number of rows after dropping NaN values: 116093

```
[6]: import glob

image_files = glob.glob("data/images/*.jpg")
num_images = len(image_files)

print(f"Number of images in 'data/images/' directory: {num_images}")
```

Number of images in 'data/images/' directory: 116162

```
[7]: import os

# Filter out the rows where the image file doesn't exist
data = data[data['uuid'].apply(lambda x: os.path.exists(f'data/images/{x}.
↪jpg'))]
```

```
[8]: task_data = data[['uuid', 'visualbrowsertimespan']]
task_data.to_csv("data/task_data.csv", index=False)
```

```
[9]: time_spans = task_data['visualbrowsertimespan'].unique()
time_span_mapping = {time_span: i for i, time_span in enumerate(time_spans)}
print(time_span_mapping)
```

```
{'1926 to 1950': 0, '1501 to 1550': 1, '1976 to 2000': 2, '1601 to 1650': 3,
'1951 to 1975': 4, '1876 to 1900': 5, '1826 to 1850': 6, '1401 to 1500': 7,
'1300 to 1400': 8, '1751 to 1775': 9, '1801 to 1825': 10, '1851 to 1875': 11,
'1901 to 1925': 12, '2001 to present': 13, '1726 to 1750': 14, '1651 to 1700':
15, '1776 to 1800': 16, '1551 to 1600': 17, '1701 to 1725': 18, 'before 1300':
19}
```

```
[10]: task_data['visualbrowsertimespan'] = task_data['visualbrowsertimespan'].
↪map(time_span_mapping)
task_data.head()
```

/tmp/ipykernel_18622/55327113.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
task_data['visualbrowsertimespan'] =
task_data['visualbrowsertimespan'].map(time_span_mapping)
```

```
[10]:
```

	uuid	visualbrowsertimespan
0	00007f61-4922-417b-8f27-893ea328206c	0
1	0000bd8c-39de-4453-b55d-5e28a9beed38	0
2	0001668a-dd1c-48e8-9267-b6d1697d43c8	0
3	00032658-8a7a-44e3-8bb8-df8c172f521d	1
4	0003d4e4-d7fd-4835-8d27-1e9e20672e39	2

```
[11]: from sklearn.model_selection import train_test_split

# Split the data into training and test sets
train_data, test_data = train_test_split(task_data, test_size=0.2,
    random_state=42)

# Print the shapes of the training and test sets
print("Training data shape:", train_data.shape)
print("Test data shape:", test_data.shape)
```

Training data shape: (92803, 2)

Test data shape: (23201, 2)

```
[12]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.tensorboard import SummaryWriter
import numpy as np
import matplotlib.pyplot as plt
import torchvision.transforms as transforms
from torch.utils.data import Dataset, DataLoader
import torchvision.transforms as transforms
from torchvision import models
from tqdm import tqdm
from PIL import Image, ImageFile
ImageFile.LOAD_TRUNCATED_IMAGES = True
from torch.nn import DataParallel

import warnings
warnings.filterwarnings("ignore")
```

```
[13]: class ImageDataset(Dataset):
    def __init__(self, data, image_dir, transform=None):
        self.data = data
        self.image_dir = image_dir
        self.transform = transform

    def __len__(self):
        return len(self.data)

    def __getitem__(self, idx):
        img_name = f"{self.image_dir}/{self.data.iloc[idx, 0]}.jpg"
        image = Image.open(img_name)
        label = self.data.iloc[idx, 1]

        if self.transform:
            image = self.transform(image)
```

```
return image, label
```

```
[14]: class ImageClassifier(nn.Module):
    def __init__(self, num_classes):
        super(ImageClassifier, self).__init__()
        self.model = models.resnet50(pretrained=True)
        in_features = self.model.fc.in_features
        self.model.fc = nn.Linear(in_features, num_classes)

    def forward(self, x):
        return self.model(x)
```

```
[15]: train_transform = transforms.Compose([
    transforms.RandomRotation(360),
    transforms.RandomHorizontalFlip(),
    transforms.RandomAffine(degrees=(-10, 10), translate=(0.05, 0.05),
↪scale=(0.95, 1.05), shear=5),
    transforms.RandomPerspective(distortion_scale=0.1, p=0.1),
    transforms.GaussianBlur(kernel_size=(3, 3), sigma=(0.1, 0.2)),
    transforms.ColorJitter(brightness=0.1, contrast=0.1, saturation=0.
↪1, hue=0.1),
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
])

test_transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

train_dataset = ImageDataset(train_data, "data/images",
↪transform=train_transform)
test_dataset = ImageDataset(test_data, "data/images", transform=test_transform)

train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = ImageClassifier(num_classes=len(time_spans)).to(device)
model = DataParallel(model)
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.0001)
summary = SummaryWriter()
iterations = 0
```

```

def train(model, train_loader, criterion, optimizer, device, epoch, summary,
↪ iterations):
    model.train()
    running_loss = 0.0
    for i, data in enumerate(tqdm(train_loader)):
        inputs, labels = data
        inputs, labels = inputs.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        summary.add_scalar("train_loss", loss.item(), iterations)
        iterations += 1
        running_loss += loss.item()
    epoch_loss = running_loss / len(train_loader)
    print(f"Train Loss: {epoch_loss}")
    return model, iterations

def test(model, test_loader, criterion, device, epoch, summary):
    model.eval()
    running_loss = 0.0
    correct = 0
    total = 0
    with torch.no_grad():
        for data in tqdm(test_loader):
            inputs, labels = data
            inputs, labels = inputs.to(device), labels.to(device)

            outputs = model(inputs)
            loss = criterion(outputs, labels)
            running_loss += loss.item()

            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    epoch_loss = running_loss / len(test_loader)
    accuracy = 100 * correct / total
    print(f"Test Loss: {epoch_loss}, Accuracy: {accuracy}")
    summary.add_scalar("test_loss", epoch_loss, epoch)
    summary.add_scalar("test_accuracy", accuracy, epoch)
    return accuracy

```

```
[16]: best_accuracy = 0
      for epoch in range(10):
          print(f"Epoch {epoch + 1}")
          model, iterations = train(model, train_loader, criterion, optimizer,
          ↪device, epoch, summary, iterations)
          accuracy = test(model, test_loader, criterion, device, epoch, summary)
          if accuracy > best_accuracy:
              best_accuracy = accuracy
              torch.save(model.state_dict(), "model.pth")
```

Epoch 1

100%| | 2901/2901 [15:52<00:00, 3.05it/s]

Train Loss: 1.690137814834421

100%| | 726/726 [01:31<00:00, 7.93it/s]

Test Loss: 1.6270704830973601, Accuracy: 49.32977026852291

Epoch 2

100%| | 2901/2901 [15:13<00:00, 3.17it/s]

Train Loss: 1.4655975448760605

100%| | 726/726 [01:29<00:00, 8.12it/s]

Test Loss: 1.6822010135026675, Accuracy: 50.592646868669455

Epoch 3

100%| | 2901/2901 [15:13<00:00, 3.17it/s]

Train Loss: 1.3722520290017743

100%| | 726/726 [01:35<00:00, 7.59it/s]

Test Loss: 1.4545212095351916, Accuracy: 54.30369380630145

Epoch 4

100%| | 2901/2901 [15:13<00:00, 3.18it/s]

Train Loss: 1.308250688450454

100%| | 726/726 [01:33<00:00, 7.79it/s]

Test Loss: 1.5210858398725178, Accuracy: 54.217490625404075

Epoch 5

100%| | 2901/2901 [15:09<00:00, 3.19it/s]

Train Loss: 1.2609052630960016

100%| | 726/726 [01:32<00:00, 7.83it/s]

Test Loss: 1.3968077205100993, Accuracy: 55.92000344812723

Epoch 6

100%| | 2901/2901 [15:20<00:00, 3.15it/s]

Train Loss: 1.2195870250062835
 100%| | 726/726 [01:32<00:00, 7.81it/s]
 Test Loss: 1.3554734793740528, Accuracy: 56.99754320934443
 Epoch 7
 100%| | 2901/2901 [15:15<00:00, 3.17it/s]
 Train Loss: 1.184016942443703
 100%| | 726/726 [01:31<00:00, 7.93it/s]
 Test Loss: 1.3625999040511685, Accuracy: 57.954398517305286
 Epoch 8
 100%| | 2901/2901 [15:11<00:00, 3.18it/s]
 Train Loss: 1.153340318035726
 100%| | 726/726 [01:30<00:00, 8.06it/s]
 Test Loss: 1.343617481141051, Accuracy: 57.967328994439896
 Epoch 9
 100%| | 2901/2901 [15:20<00:00, 3.15it/s]
 Train Loss: 1.1244382328388816
 100%| | 726/726 [01:31<00:00, 7.90it/s]
 Test Loss: 1.279041744132344, Accuracy: 59.18710400413775
 Epoch 10
 100%| | 2901/2901 [15:12<00:00, 3.18it/s]
 Train Loss: 1.0983684018475153
 100%| | 726/726 [01:31<00:00, 7.89it/s]
 Test Loss: 1.3050892327606842, Accuracy: 59.29485798025947

```
[18]: for epoch in range(30):
      print(f"Epoch {epoch + 1}")
      model, iterations = train(model, train_loader, criterion, optimizer,
      ↪device, epoch, summary, iterations)
      accuracy = test(model, test_loader, criterion, device, epoch, summary)
      if accuracy > best_accuracy:
          best_accuracy = accuracy
          torch.save(model.state_dict(), "model.pth")
```

Epoch 1
 100%| | 2901/2901 [15:17<00:00, 3.16it/s]
 Train Loss: 1.0731199352402638
 100%| | 726/726 [01:33<00:00, 7.79it/s]

Test Loss: 1.3640705223911065, Accuracy: 58.84660143959312
Epoch 2

100%| | 2901/2901 [15:11<00:00, 3.18it/s]

Train Loss: 1.0487641955564369

100%| | 726/726 [01:31<00:00, 7.97it/s]

Test Loss: 1.360556701983302, Accuracy: 58.600922374035605
Epoch 3

100%| | 2901/2901 [15:11<00:00, 3.18it/s]

Train Loss: 1.026329971066593

100%| | 726/726 [01:30<00:00, 8.04it/s]

Test Loss: 1.3139927555132176, Accuracy: 59.45864402396448
Epoch 4

100%| | 2901/2901 [15:14<00:00, 3.17it/s]

Train Loss: 1.0068358256668108

100%| | 726/726 [01:31<00:00, 7.95it/s]

Test Loss: 1.2765662414804306, Accuracy: 60.82927460023275
Epoch 5

100%| | 2901/2901 [15:17<00:00, 3.16it/s]

Train Loss: 0.9853862111464076

100%| | 726/726 [01:31<00:00, 7.93it/s]

Test Loss: 1.3261786331486767, Accuracy: 60.4629110814189
Epoch 6

100%| | 2901/2901 [15:14<00:00, 3.17it/s]

Train Loss: 0.9690948955480989

100%| | 726/726 [01:31<00:00, 7.98it/s]

Test Loss: 1.311826701458164, Accuracy: 60.863755872591696
Epoch 7

100%| | 2901/2901 [15:13<00:00, 3.17it/s]

Train Loss: 0.9511676182709082

100%| | 726/726 [01:31<00:00, 7.92it/s]

Test Loss: 1.3074272146484411, Accuracy: 61.4154562303349
Epoch 8

100%| | 2901/2901 [15:11<00:00, 3.18it/s]

Train Loss: 0.9290971909256402

100%| | 726/726 [01:30<00:00, 8.01it/s]
Test Loss: 1.2541883173231907, Accuracy: 61.56200163786044
Epoch 9
100%| | 2901/2901 [15:17<00:00, 3.16it/s]
Train Loss: 0.9111072075588873
100%| | 726/726 [01:31<00:00, 7.92it/s]
Test Loss: 1.3668394812867661, Accuracy: 60.949959053489074
Epoch 10
100%| | 2901/2901 [15:09<00:00, 3.19it/s]
Train Loss: 0.8960632865735475
100%| | 726/726 [01:32<00:00, 7.85it/s]
Test Loss: 1.3644437264066098, Accuracy: 60.678419033662344
Epoch 11
100%| | 2901/2901 [15:15<00:00, 3.17it/s]
Train Loss: 0.8810253038832913
100%| | 726/726 [01:33<00:00, 7.76it/s]
Test Loss: 1.3553913240590372, Accuracy: 60.71721046506616
Epoch 12
100%| | 2901/2901 [15:20<00:00, 3.15it/s]
Train Loss: 0.8620052225155815
100%| | 726/726 [01:30<00:00, 8.01it/s]
Test Loss: 1.287482999244669, Accuracy: 61.958536269988365
Epoch 13
100%| | 2901/2901 [15:09<00:00, 3.19it/s]
Train Loss: 0.845609387511921
100%| | 726/726 [01:32<00:00, 7.83it/s]
Test Loss: 1.4490600350228222, Accuracy: 60.35946726434205
Epoch 14
100%| | 2901/2901 [15:15<00:00, 3.17it/s]
Train Loss: 0.8290279850348322
100%| | 726/726 [01:36<00:00, 7.53it/s]
Test Loss: 1.3484457275591606, Accuracy: 62.26886772121891
Epoch 15
100%| | 2901/2901 [15:12<00:00, 3.18it/s]

Train Loss: 0.8145340183548827
100%| | 726/726 [01:34<00:00, 7.67it/s]
Test Loss: 1.294868934909831, Accuracy: 63.130899530192664
Epoch 16
100%| | 2901/2901 [15:13<00:00, 3.18it/s]
Train Loss: 0.7965203226639952
100%| | 726/726 [01:33<00:00, 7.78it/s]
Test Loss: 1.3278530860324864, Accuracy: 62.29472867548812
Epoch 17
100%| | 2901/2901 [15:15<00:00, 3.17it/s]
Train Loss: 0.7809749720056317
100%| | 726/726 [01:31<00:00, 7.94it/s]
Test Loss: 1.3749777219019645, Accuracy: 62.16542390414206
Epoch 18
100%| | 2901/2901 [15:11<00:00, 3.18it/s]
Train Loss: 0.7681955712411125
100%| | 726/726 [01:31<00:00, 7.91it/s]
Test Loss: 1.375300587044603, Accuracy: 61.81630102150769
Epoch 19
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.7506349904501861
100%| | 726/726 [01:31<00:00, 7.90it/s]
Test Loss: 1.4458212725428181, Accuracy: 63.393819231929655
Epoch 20
100%| | 2901/2901 [15:13<00:00, 3.18it/s]
Train Loss: 0.7362199015388896
100%| | 726/726 [01:32<00:00, 7.86it/s]
Test Loss: 1.4028292325820477, Accuracy: 62.863669669410804
Epoch 21
100%| | 2901/2901 [15:10<00:00, 3.18it/s]
Train Loss: 0.7200810335825657
100%| | 726/726 [01:33<00:00, 7.77it/s]
Test Loss: 1.426847750490362, Accuracy: 62.69557346666092
Epoch 22

100%| | 2901/2901 [15:16<00:00, 3.17it/s]
Train Loss: 0.7063657639458935
100%| | 726/726 [01:35<00:00, 7.63it/s]
Test Loss: 1.4502108799031943, Accuracy: 62.28610835739839
Epoch 23
100%| | 2901/2901 [15:13<00:00, 3.17it/s]
Train Loss: 0.6888736269517999
100%| | 726/726 [01:32<00:00, 7.85it/s]
Test Loss: 1.4406594608194572, Accuracy: 62.79901728373777
Epoch 24
100%| | 2901/2901 [15:14<00:00, 3.17it/s]
Train Loss: 0.6791201755430155
100%| | 726/726 [01:31<00:00, 7.95it/s]
Test Loss: 1.5414045637591482, Accuracy: 62.53178742295591
Epoch 25
100%| | 2901/2901 [15:17<00:00, 3.16it/s]
Train Loss: 0.66138602537416
100%| | 726/726 [01:30<00:00, 7.98it/s]
Test Loss: 1.4996483213891667, Accuracy: 63.0015947588466
Epoch 26
100%| | 2901/2901 [15:16<00:00, 3.17it/s]
Train Loss: 0.6498614293014292
100%| | 726/726 [01:31<00:00, 7.98it/s]
Test Loss: 1.5028258647882577, Accuracy: 62.863669669410804
Epoch 27
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.6370408145344204
100%| | 726/726 [01:31<00:00, 7.97it/s]
Test Loss: 1.5096620853897953, Accuracy: 62.72143442093013
Epoch 28
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.623176160859511
100%| | 726/726 [01:30<00:00, 8.05it/s]

Test Loss: 1.5724494867745182, Accuracy: 62.54471790009051
Epoch 29

100%| | 2901/2901 [15:13<00:00, 3.18it/s]

Train Loss: 0.6076582055160071

100%| | 726/726 [01:30<00:00, 8.02it/s]

Test Loss: 1.641739511210728, Accuracy: 61.9671565880781
Epoch 30

100%| | 2901/2901 [15:11<00:00, 3.18it/s]

Train Loss: 0.5941098445563512

100%| | 726/726 [01:31<00:00, 7.92it/s]

Test Loss: 1.5748681583680397, Accuracy: 62.574889013404594

```
[19]: for epoch in range(30):  
        print(f"Epoch {epoch + 1}")  
        model, iterations = train(model, train_loader, criterion, optimizer,   
        ↪device, epoch, summary, iterations)  
        accuracy = test(model, test_loader, criterion, device, epoch, summary)  
        if accuracy > best_accuracy:  
            best_accuracy = accuracy  
            torch.save(model.state_dict(), "model.pth")
```

Epoch 1

100%| | 2901/2901 [15:13<00:00, 3.18it/s]

Train Loss: 0.5841842094029858

100%| | 726/726 [01:31<00:00, 7.94it/s]

Test Loss: 1.574741728583315, Accuracy: 63.10934873496832
Epoch 2

100%| | 2901/2901 [15:17<00:00, 3.16it/s]

Train Loss: 0.5704419449952831

100%| | 726/726 [01:31<00:00, 7.93it/s]

Test Loss: 1.6124020591172963, Accuracy: 63.21710271109004
Epoch 3

100%| | 2901/2901 [15:12<00:00, 3.18it/s]

Train Loss: 0.5575335187897194

100%| | 726/726 [01:30<00:00, 8.03it/s]

Test Loss: 1.6218958239253232, Accuracy: 63.46709193569243
Epoch 4
100%| | 2901/2901 [15:07<00:00, 3.20it/s]
Train Loss: 0.5435561933925505
100%| | 726/726 [01:30<00:00, 8.02it/s]
Test Loss: 1.6054723077240727, Accuracy: 63.2774449377182
Epoch 5
100%| | 2901/2901 [15:30<00:00, 3.12it/s]
Train Loss: 0.5321687847550512
100%| | 726/726 [01:36<00:00, 7.55it/s]
Test Loss: 1.654402204288924, Accuracy: 62.587819490539204
Epoch 6
100%| | 2901/2901 [15:15<00:00, 3.17it/s]
Train Loss: 0.5229305098207603
100%| | 726/726 [01:31<00:00, 7.93it/s]
Test Loss: 1.7675459911487483, Accuracy: 62.574889013404594
Epoch 7
100%| | 2901/2901 [15:27<00:00, 3.13it/s]
Train Loss: 0.5102772527971131
100%| | 726/726 [01:33<00:00, 7.79it/s]
Test Loss: 1.75020331283903, Accuracy: 62.80332744278264
Epoch 8
100%| | 2901/2901 [15:13<00:00, 3.18it/s]
Train Loss: 0.5003935639373766
100%| | 726/726 [01:31<00:00, 7.92it/s]
Test Loss: 1.7361015405290383, Accuracy: 62.850739192276194
Epoch 9
100%| | 2901/2901 [15:17<00:00, 3.16it/s]
Train Loss: 0.48540374629972555
100%| | 726/726 [01:30<00:00, 7.99it/s]
Test Loss: 1.6879902401164215, Accuracy: 63.316236369122024
Epoch 10
100%| | 2901/2901 [15:11<00:00, 3.18it/s]
Train Loss: 0.4764779852186347

100%| | 726/726 [01:31<00:00, 7.90it/s]
Test Loss: 1.764539520230565, Accuracy: 62.2430067669497
Epoch 11
100%| | 2901/2901 [15:11<00:00, 3.18it/s]
Train Loss: 0.46535194102226557
100%| | 726/726 [01:33<00:00, 7.73it/s]
Test Loss: 1.729557144006627, Accuracy: 62.967113486487655
Epoch 12
100%| | 2901/2901 [15:14<00:00, 3.17it/s]
Train Loss: 0.45684047246217646
100%| | 726/726 [01:29<00:00, 8.07it/s]
Test Loss: 1.8249542707007778, Accuracy: 62.36369122020603
Epoch 13
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.44226976285275404
100%| | 726/726 [01:31<00:00, 7.95it/s]
Test Loss: 1.6902453749550932, Accuracy: 63.65242877462178
Epoch 14
100%| | 2901/2901 [15:10<00:00, 3.19it/s]
Train Loss: 0.4369557541311917
100%| | 726/726 [01:30<00:00, 8.01it/s]
Test Loss: 1.838951381441647, Accuracy: 61.83354165768717
Epoch 15
100%| | 2901/2901 [15:18<00:00, 3.16it/s]
Train Loss: 0.43023958494543246
100%| | 726/726 [01:32<00:00, 7.89it/s]
Test Loss: 1.8583025800130912, Accuracy: 62.52316710486617
Epoch 16
100%| | 2901/2901 [15:13<00:00, 3.17it/s]
Train Loss: 0.4164440942474251
100%| | 726/726 [01:31<00:00, 7.97it/s]
Test Loss: 1.8575779050999108, Accuracy: 63.26020430153873
Epoch 17
100%| | 2901/2901 [15:21<00:00, 3.15it/s]

Train Loss: 0.4069580955675494
100%| | 726/726 [01:29<00:00, 8.07it/s]
Test Loss: 1.9286355537570212, Accuracy: 62.337830265936816
Epoch 18
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.3958651038237951
100%| | 726/726 [01:32<00:00, 7.81it/s]
Test Loss: 1.9415915090354678, Accuracy: 62.842118874186454
Epoch 19
100%| | 2901/2901 [15:13<00:00, 3.18it/s]
Train Loss: 0.38871323414498754
100%| | 726/726 [01:36<00:00, 7.56it/s]
Test Loss: 1.9688646134474, Accuracy: 63.00590491789147
Epoch 20
100%| | 2901/2901 [15:16<00:00, 3.16it/s]
Train Loss: 0.37920100117209493
100%| | 726/726 [01:30<00:00, 8.04it/s]
Test Loss: 2.0810656687116493, Accuracy: 61.65251497780268
Epoch 21
100%| | 2901/2901 [15:14<00:00, 3.17it/s]
Train Loss: 0.37289904823583886
100%| | 726/726 [01:31<00:00, 7.91it/s]
Test Loss: 1.855735930775808, Accuracy: 62.5188569458213
Epoch 22
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.3641731788329485
100%| | 726/726 [01:32<00:00, 7.88it/s]
Test Loss: 2.043836662920382, Accuracy: 62.37231153829576
Epoch 23
100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.35301943851319606
100%| | 726/726 [01:30<00:00, 8.00it/s]
Test Loss: 2.0979382489133145, Accuracy: 63.75587259169863
Epoch 24

100%| | 2901/2901 [15:12<00:00, 3.18it/s]
Train Loss: 0.348063370968629
100%| | 726/726 [01:30<00:00, 8.02it/s]
Test Loss: 1.9430203689852366, Accuracy: 64.07482436101893
Epoch 25
100%| | 2901/2901 [15:15<00:00, 3.17it/s]
Train Loss: 0.3407567353576925
100%| | 726/726 [01:30<00:00, 8.06it/s]
Test Loss: 2.115215538726406, Accuracy: 62.35507090211629
Epoch 26
100%| | 2901/2901 [15:13<00:00, 3.18it/s]
Train Loss: 0.33419982180184266
100%| | 726/726 [01:30<00:00, 8.04it/s]
Test Loss: 2.039065026547298, Accuracy: 63.333477005301496
Epoch 27
100%| | 2901/2901 [15:10<00:00, 3.19it/s]
Train Loss: 0.32518909942976404
100%| | 726/726 [01:31<00:00, 7.98it/s]
Test Loss: 2.0296253873132803, Accuracy: 63.24727382440412
Epoch 28
100%| | 2901/2901 [15:13<00:00, 3.18it/s]
Train Loss: 0.3185634356260464
100%| | 726/726 [01:30<00:00, 8.04it/s]
Test Loss: 2.1329430344183584, Accuracy: 62.09215120037929
Epoch 29
100%| | 2901/2901 [15:09<00:00, 3.19it/s]
Train Loss: 0.3135327958443768
100%| | 726/726 [01:33<00:00, 7.72it/s]
Test Loss: 2.167870124591612, Accuracy: 61.93698547476402
Epoch 30
100%| | 2901/2901 [15:16<00:00, 3.16it/s]
Train Loss: 0.30488706763316253
100%| | 726/726 [01:33<00:00, 7.74it/s]
Test Loss: 2.082339142974833, Accuracy: 63.59639670703849

```
[20]: # load best model and check accuracy
model.load_state_dict(torch.load("model.pth"))
test(model, test_loader, criterion, device, 0, summary)

100%|      | 726/726 [01:34<00:00, 7.70it/s]

Test Loss: 1.9430203689852366, Accuracy: 64.07482436101893
```

```
[20]: 64.07482436101893
```

```
[ ]:
```